

Unit 12: Shortest Path Problems

Agenda:


- ▶ Single-Source Shortest Path problem (SSSP)
 - ▶ Dijkstra's algorithm for the non-negative weights case
 - ▶ Bellman-Ford algorithm for graphs with no negative cycles
- ▶ All-Pairs Shortest Path problem (APSP)
 - ▶ Floyd-Washall algorithm

Reading:

- ▶ CLRS : 643-650, 658-664, 651-655, 684-686, 693-699

Shortest path problems:

- ▶ BFS recall: outputs every s -to- v shortest path
 - ▶ s — start vertex
 - ▶ v — reachable vertex from s (residing in a same connected component)
 - ▶ shortest — # edges
 - ▶ running time $\Theta(n + m)$
 - ▶ But what if the edges of the graph have weights?
- In this case, shortest-path in terms of #edges isn't shortest weighted path.




E.g. shortest weighted-path distance between 1 and 2 is 8.

- ▶ The weight of a path = sum of weights on edges on path.
Note: if there is no path between two nodes, the distance is set to ∞ ...
- ▶ The SHORTEST PATH problem: find the shortest path in an edge-weighted graph connecting s and t .
- ▶ Turns out, shortest-path has a few properties that allow us to infer in the process **all** shortest-paths from a node s to any other node in the graph.
So we study the SINGLE-SOURCE SHORTEST PATH (SSSP) problem:
Given an edge-weighted graph G and a source s , find out for each vertex $v \in V(G)$ a shortest paths from s to v .
- ▶ Variants:
 - ▶ edge weights: non-negative vs. arbitrary weights

Shortest Paths

- ▶ A shortest path from u to v : out of all $u \rightarrow v$ paths, it is a path of minimal weight. (Can be more than one.)
- ▶ But a shortest path must satisfy **subpath optimality**:
if $(u_0, u_1, \dots, \underbrace{u_i, \dots, u_j}_{\text{optimal}}, \dots, u_k)$ is a shortest $u_0 \rightarrow u_k$ path, then $(u_i, u_{i+1}, \dots, u_j)$ is a shortest $u_i \rightarrow u_j$ path.
- ▶ we denote $d(u, v)$ as the shortest weighted path length from u to v
- ▶ Shortest paths remain well-defined if there are negative weight directed



edges... (why not undirected?)

- ▶ ...as long as there isn't a negative weight path from a node to itself.
- ▶ So we assume **no negative cycles**.
- ▶ Shortest path distances — d is a metric:
 - ▶ For any u , we have $d(u, u) = 0$
 - ▶ For any u, v , we have $d(u, v) = d(v, u)$ (if G is undirected)
 - ▶ For any u, v, w , we have $d(u, w) \leq d(u, v) + d(v, w)$.

Common Outline of Single Source Shortest Path Algorithms

- ▶ Our shortest paths algorithm starts at a source s .
- ▶ It maintains a *dist* attribute for each vertex that will serve as the estimation of the shortest-path distance.
- ▶ During the execution of the algorithm, we **always** have $u.dist \geq d(s, u)$.
- ▶ We start with `init()`: set $s.dist \leftarrow 0$ and $u.dist \leftarrow \infty$ for any $u \neq s$.
- ▶ We only update the *dist* attribute by Relaxing

```

procedure relax( $u, v$ )           ** $u, v$  two adjacent nodes
  if ( $v.dist > u.dist + w(u, v)$ ) then
     $v.dist \leftarrow u.dist + w(u, v)$ 

```

- ▶ Claim: any algorithm that starts with `init()` and only updates *dist* using `relax()` must always satisfy $v.dist \geq d(s, v)$ for any v .
- ▶ Proof: by induction on the number of times `relax()` is invoked.
 Base case: invoked 0 times — claim holds through `init()`.
 Induction step: If `relax(u, v)` doesn't change $v.dist$ we are done.
 Otherwise, we now have

$$v.dist = u.dist + w(u, v) \geq d(s, u) + w(u, v) \geq d(s, v)$$

- ▶ Corollary: Since `relax()` cannot increase $v.dist$, so if and when we set $v.dist = d(s, v)$ we keep $v.dist$ unchanged.

Dijkstra's SSSP algorithm:

- ▶ For graphs with non-negative weights (both directed and undirected)
- ▶ Idea in Dijkstra's algorithm:
 - ▶ Maintains a set S of vertices for which we know the shortest path. Initially, $S = \{s\}$, at the end: $S = V$.
 - ▶ Which vertex from $\bar{S} = V \setminus S$ should we pick?
 - ▶ Dijkstra: the greedy solution — the node in \bar{S} with minimum *dist*.

▶ procedure dijkstra(G, w, s) **** $G = (V, E)$, w = weights**
 foreach v do ****initialization**

$v.dist \leftarrow \infty$

$v.predec \leftarrow \text{NIL}$

$s.dist \leftarrow 0$

 Build Min-Priority-Queue Q on all nodes, key = *dist*

**** nodes in Q are nodes we are not yet sure about**

**** namely Q holds nodes in \bar{S}**

 while ($Q \neq \emptyset$) do

$u \leftarrow \text{ExtractMin}(Q)$ **** s dequeued first**

 foreach v neighbor of u do

 if ($v.dist > u.dist + w(u, v)$) then


$v.dist \leftarrow u.dist + w(u, v)$ ****a Relax() call**

$v.predec \leftarrow u$


 decrease-key($Q, v, v.dist$)

Dijkstra's SSSP algorithm — an example:

- Input graph G :



- $\text{dijkstra}(G, 1)$:



- $\text{dijkstra}(G, 1)$ trace:

v	1	2	3	4	5	6	7
$v.\text{dist}/v.\text{predec}$	0/NIL	∞ /NIL	∞ /NIL	∞ /NIL	∞ /NIL	∞ /NIL	∞ /NIL
1 dequeued	0/NIL	9/1	∞ /NIL	∞ /NIL	3/1	12/1	∞ /NIL
5 dequeued	0/NIL	9/1	11/5	∞ /NIL	3/1	12/1	∞ /NIL
2 dequeued	0/NIL	9/1	10/2	∞ /NIL	3/1	12/1	19/2
3 dequeued	0/NIL	9/1	10/2	25/3	3/1	12/1	19/2
6 dequeued	0/NIL	9/1	10/2	25/3	3/1	12/1	19/2
7 dequeued	0/NIL	9/1	10/2	23/7	3/1	12/1	19/2

Dijkstra's SSSP algorithm — Correctness:

- ▶ LI: “at the start of each while-loop iteration $u.dist = d(s, u)$ for all $u \notin Q$.”
- ▶ Initialization: The statement is vacuously true when Q holds all nodes. The statement is clearly true for the first vertex taken out of Q : the source s .
- ▶ Termination: At the end of the while-loop, Q is empty, so we have found all shortest-paths distances from s .
- ▶ Maintenance:
 - ▶ Suppose LI holds at the beginning of the iteration. Denote u as the node $\text{Extract-Min}(Q)$ takes out.
 - ▶ ASOC $d(s, u) < u.dist$.
 - ▶ Look at a shortest-path $s \rightarrow u$. $s \notin Q$ but $u \in Q$. Let (x, y) be the first edge such that $x \notin Q$ but $y \in Q$. (y might be u , x might be s .)
 - ▶ First, $d(s, y) \leq d(s, u)$ (subpath optimality + non negative weights)
 - ▶ Second, when we took x out of Q we made sure $y.dist \leq x.dist + w(x, y)$
 - ▶ Since $x \notin Q$ in the beginning of the iteration, then $x.dist = d(s, x)$.
 - ▶ Thus $y.dist \leq d(s, x) + w(x, y) = d(s, y)$ (Again, subpath optimality)
 - ▶ Altogether: $y.dist = d(s, y) \leq d(s, u) < u.dist$
 - ▶ If $y = u$ — immediate contradiction.
If $y \neq u$ — then $\text{Extract-Min}(Q)$ should return y not u .
Contradiction in any case.

Dijkstra's SSSP algorithm — analysis:

- ▶ $|V| = n$ and $|E| = m$
- ▶ Running time:
 - ▶ `init()` — takes $O(n)$ time.
 - ▶ Initializing the priority-queue — $O(n)$
 - ▶ For each node, `ExtractMin` takes $O(\log(n))$ time.
 - ▶ ... and we search for all neighbors ($O(\deg(v))$ in the adjacency list model, $O(n)$ in the adjacency matrix model)
 - ▶ For every edge, we examine the edge at most twice (each time we take its endpoints from Q) and invoke `relax()` at most once, and decrease the key at most once.
So $O(\log(n))$ work per edge.

- ▶ In the adjacency-list model

$$O(n) + O(n) + O(n \log(n)) + O\left(\sum_v \deg(v)\right) + O(m \log(n)) = O((n+m) \log(n))$$

- ▶ In the adjacency-matrix model

$$O(n) + O(n) + O(n(\log(n) + n)) + O(m \log(n)) = O(n^2 + m \log(n))$$

- ▶ There exists a more refined implementation of the PQ (with Fibonacci heaps) that gives runtime $O(n \log(n) + m)$

Dealing with Negative Weights

- ▶ Consider a $s \rightarrow v_k$ shortest path: $(s, v_1, v_2, \dots, v_k)$ (has k edges)
- ▶ Subpath optimality: this is also a shortest path for any $s \rightarrow v_i$.
- ▶ Suppose that among the `relax()` calls of our algorithm it also made a subsequence of calls (in this order):
`relax(s, v1), relax(v1, v2), relax(v2, v3), ... , relax(vk-1, vk),`
- ▶ The first call sets
 $v_1.dist \leq 0 + w(s, v_1) = d(s, v_1) \Rightarrow v_1.dist = d(s, v_1)$
- ▶ The second call sets
 $v_2.dist \leq v_1.dist + w(v_1, v_2) = d(s, v_2) \Rightarrow v_2.dist = d(s, v_2)$
- ▶ The third call sets
 $v_3.dist \leq v_2.dist + w(v_2, v_3) = d(s, v_3) \Rightarrow v_3.dist = d(s, v_3)$
- ▶ ...the k -th call sets $v_k.dist = d(s, v_k)$
- ▶ How can we make sure we have such k calls?
- ▶ Brute force solution: makes all the `relax()` calls on all edges k times.
- ▶ What's k ? At most $n - 1$...

Bellman-Ford Algorithm

- ▶ procedure Bellman-Ford(G, s)


```

foreach  $v \in V(G)$  do
     $v.dist \leftarrow \infty$                                 ** Initialization
     $v.predec \leftarrow \text{NIL}$ 
 $s.dist \leftarrow 0$ 
for ( $i$  from 1 to  $n - 1$ ) do
    foreach edge  $(u, v)$  do
        if ( $v.dist > u.dist + w(u, v)$ ) then
             $v.dist \leftarrow u.dist + w(u, v)$           ** Relax()
             $v.predec \leftarrow u$ 
foreach edge  $(u, v)$  do
    if ( $v.dist > u.dist + w(u, v)$ ) then
        return ‘‘negative-cycle’’


```
- ▶ Why is the negative cycle-check true? (last 3 lines)
 - ▶ Along a negative cycle, we keep relaxing and relaxing infinitely many times...
- ▶ Runtime on a graph with n nodes and m edges:
 - $O(n \cdot m)$ in the adjacency list model
 - $O(n^2 + n \cdot m)$ in the adjacency matrix model (first construct the array of edges in $O(n^2)$ time)

Bellman-Ford algorithm — an example:

- Input graph G :



- Bellman-Ford($G, 3$):




- Bellman-Ford($G, 3$) trace:

Edge order: $(4, 2), (1, 2), (1, 4), (3, 1), (3, 4)$


v	1	2	3	4
$v.dist/v.predec$	∞/NIL	∞/NIL	0/NIL	∞/NIL
round 1	-3/3	∞/NIL	0/NIL	16/3
round 2	-3/3	6/1	0/NIL	-15/1
round 3	-3/3	-8/4	0/NIL	-15/1

Bellman-Ford algorithm — an example with negative cycle:

- Input graph G :



- Bellman-Ford($G, 3$):



- Bellman-Ford($G, 3$) trace:

Edge order: $(4, 2), (2, 1), (1, 4), (3, 1), (3, 4)$

v $v.dist/v.predec$	1	2	3	4
	∞/NIL	∞/NIL	0/NIL	∞/NIL
round 1	-3/3	∞/NIL	0/NIL	16/NIL
round 2	-3/3	23/4	0/NIL	-15/1
round 3	-3/3	-8/4	0/NIL	-15/1
checkup round	-4/2	-8/4	0/NIL	-15/1

All-Pairs Shortest-Path Problem:

- ▶ General case (weights can be negative but no negative cycle);
- ▶ Output: shortest path between every pair u, v of vertices.
- ▶ If the graph has no negative weights: run Dijkstra n times with each vertex as a source.
- ▶ Runtime:
 - ▶ $O(n(n + m) \log(n))$.
 - ▶ ($O(n^2 \log(n) + nm)$ with the improved implementation that uses a Fibonacci-heap.)
- ▶ If the graph has negative weights: run Bellman-Ford n times with each vertex as a source.
 - ▶ $O(n^2m)$.
- ▶ Can we do better?

Floyd-Warshall's algorithm for All-Pairs-Shortest path:

- ▶ General case (weights can be negative but no negative cycle);
- ▶ Idea: Use dynamic programming.
 Define $d[i, j, k]$ to be the length of shortest path from i to j for which all intermediate vertices are in $\{1, \dots, k\}$, for every $1 \leq i, j \leq n$ and $k \leq n$.
- ▶ When $k = 0 \implies$ no intermediate vertex, so: $d[i, j, 0] = w(i, j)$ if the edge (i, j) exists, ∞ otherwise.
- ▶ For general $k \geq 1$:
 - ▶ If the path does not contain k , inter. vertices only from $\{1, \dots, k - 1\}$: $d[i, j, k - 1]$.
 - ▶ If the path contains k , it has two parts: one goes from i to k with intermediate only from $\{1, \dots, k - 1\}$, followed by a path from k to j using only from $\{1, \dots, k - 1\}$: $d[i, k, k - 1] + d[k, j, k - 1]$.
- ▶ Recurrence for $k \geq 1$:

$$d[i, j, k] = \min \begin{cases} d[i, j, k - 1] & \text{path doesn't use vertex } k \\ d[i, k, k - 1] + d[k, j, k - 1] & \text{path does use vertex } k \end{cases}$$
- ▶ We compute the table bottom-up, starting from smaller values of k to larger values.

Floyd-Warshall's algorithm:

- ▶ procedure Floyd-Warshall(G)
 - for (i from 1 to n) do
 - for (j from 1 to n) do
 - if ((i, j) is an edge) do
 - $d[i, j, 0] \leftarrow w(i, j)$
 - $b[i, j] \leftarrow 0$ ** $b[i, j]$ keeps the break point vertex
 - else
 - $d[i, j, 0] \leftarrow \infty$
 - $b[i, j] \leftarrow \perp$ ** \perp stands for “don’t know”
 - for (k from 1 to n) do
 - for (i from 1 to n) do
 - for (j from 1 to n) do
 - if ($d[i, j, k-1] > d[i, k, k-1] + d[k, j, k-1]$) then
 - $d[i, j, k] \leftarrow d[i, k, k-1] + d[k, j, k-1]$
 - $b[i, j] \leftarrow k$
 - else
 - $d[i, j, k] \leftarrow d[i, j, k-1]$
- ▶ Running time: $\Theta(n^3)$

Floyd-Warshall's algorithm:

- ▶ To print the actual path between i and j

```
procedure Print-FW( $i, j$ )
```

```
  if ( $b[i, j] = 0$ ) then
```

```
    Print “( $i, j$ )”
```

```
  else if ( $b[i, j] = \perp$ ) then
```

```
    Print “no path”
```

```
  else
```

```
    Print-FW( $i, b[i, j]$ )    ** First print the path  $i \rightarrow b[i, j]$ 
```

```
    Print-FW( $b[i, j], j$ )    ** then print the path  $b[i, j] \rightarrow j$ 
```

- ▶ Correctness follows from the correctness of the Floyd-Warshall algorithm (the fact that the matrix b doesn't contain loops).
- ▶ Running time:
 - ▶ The recurrence relation:
 If the shortest-path between i and j is of length 1 then $T(1) = O(1)$.
 If the shortest-path between i and j is of length $k > 1$ then we invoke 2 recursions: to print the path from i to $b[i, j]$ and then to print the path from $b[i, j]$ to j . Denoting the length of the path from i to $b[i, j]$ as ℓ we have that

$$T(k) = O(1) + O(\ell) + O(k - \ell)$$

- ▶ You can guess-and-test and see this solves to $\Theta(k)$.
 As $k \leq n - 1$ the runtime is $O(n)$.
- ▶ Runtime of printing all path of all $\binom{n}{2}$ possible i and j s: $O(n^3)$.